

NCUSCC 2024秋季考核——Pytorch试题实验报告


作者：杨许玮

开始时间：24/10/13

完成时间：24/10/21

实验前准备

- 实验中有要求：报告必须采用 LaTeX 或 Markdown 格式撰写。
那我们不妨准备环境能够边实验边写报告。
 - 步骤：于[官网](#)中下载vs code，下载完成后我们可以安装Markdown All in One扩展来以Markdown 格式撰写。
 - 好处：通过该扩展，我们能够实时预览来检查我们的报告撰写方便修改。
- 选择我们实验的虚拟环境
 - 选择：通过WSL2来搭建虚拟环境
 - 原因：通过VM安装的系统的硬件都为虚拟化的硬件，我们难以直连我们的显卡。故使用WSL2来搭建虚拟环境（Ubuntu 22.04 LTS 操作系统），以此我们能够实现win10和Linux可以无缝切换，十分方便，省去很多不必要的麻烦。并且WSL2搭建虚拟环境也非常的便捷快速。
- 利用vscode连接WSL便于编辑python代码
 - 步骤：下载Remote Development扩展（该扩展为几个扩展打包在一起的，它包含了Remote-WSL，Remote-SSH，Remote-Container），此时左边为出现个屏幕的图案，我们可以以此连接WSL，如图：



- 好处：使代码调试更便捷美观。

实验环境的搭建

以WSL2搭建虚拟环境

1. 启用适用于 Linux 的 Windows 子系统

需要先启用“适用于 Linux 的 Windows 子系统”可选功能，然后才能在 Windows 上安装 Linux 分发。以管理员身份打开 PowerShell (“开始”菜单 > “PowerShell” > 单击右键 > “以管理员身份运行”)，然后输入以下命令：

```
dism.exe /online /enable-feature /featurename:Microsoft-Windows-Subsystem-Linux /all /norestart
```

2. 启用虚拟机功能

安装 WSL 2 之前，必须启用“虚拟机平台”可选功能。计算机需要虚拟化功能才能使用此功能。以管理员身份打开 PowerShell 并运行：

```
dism.exe /online /enable-feature /featurename:VirtualMachinePlatform /all /norestart
```

重新启动计算机，以完成 WSL 安装并更新到 WSL 2。

3. 下载 Linux 内核更新包

首先我们需要确定自己计算机的类型， 如果不确定自己计算机的类型，请打开命令提示符或 PowerShell，并输入：

```
systeminfo | find "System Type".
```

- 如果是x64计算机，则下载[适用于 x64 计算机的 WSL2 Linux 内核更新包](#)
- 如果是ARM64计算机，则下载[ARM64包](#)

然后运行上一步中下载的更新包。

4. 将 WSL 2 设置为默认版本

打开 PowerShell，然后在安装新的 Linux 发行版时运行以下命令，将 WSL 2 设置为默认版本：

```
wsl --set-default-version 2
```

5. 安装所选的 Linux 分发

根据要求我们需要安装Ubuntu 22.04 LTS，我们可以打开 Microsoft Store在分发版的页面中，选择“获取”。

首次启动新安装的 Linux 分发版时，将打开一个控制台窗口，系统会要求你等待一分钟或两分钟，以便文件解压缩并存储到电脑上。未来的所有启动时间应不到一秒。

然后，需要为新的 Linux 分发版创建用户帐户和密码。

到此我们便完成了使用WSL2安装Ubuntu 22.04 LTS。

安装NVIDIA 驱动和 CUDA Toolkit

NVIDIA 驱动的安装

由于windows会自动为wsl安装nvidia驱动。

我们也可以通过在命令行中输入 nvidia-smi 查询自己的显卡驱动版本，如图：

```
caicai@DESKTOP-26UFLTD: ~$ nvidia-smi
Mon Oct 14 16:21:20 2024
```

NVIDIA-SMI 530.56			Driver Version: 532.10			CUDA Version: 12.1		
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M. MIG	M.
0	NVIDIA GeForce RTX 4060 L...	On	00000000:01:00.0 On				N/A	
N/A	52C	P3	628MiB / 8188MiB		12%	Default	N/A	

Processes:								
GPU	GI	CI	PID	Type	Process name		GPU Memory	Usage
	ID	ID						
No running processes found								

CUDA Toolkit的安装

可直接去[官网](#)下载所需版本，版本的选择上建议选较早的版本，因为较早的版本较为稳定bug较少，这里我选择安装11.8版本的。

Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown. By downloading and using the software, you agree to fully comply with the terms and conditions of the [CUDA EULA](#).

Operating System

LinuxWindows

Architecture

x86_64ppc64learm64-sbsaaarch64-jetson

Distribution

CentOSDebianFedoraKylinOSOpenSUSERHELRockySLESUbuntu

Version

WSL-Ubuntu2.0

Installer Type

deb (local)deb (network)runfile (local)

Download Installer for Linux WSL-Ubuntu 2.0 x86_64

The base installer is available for download below.

Base Installer

Installation Instructions:

```
$ wget https://developer.download.nvidia.com/compute/cuda/repos/wsl-ubuntu/x86_64/cuda-wsl-ubuntu.pin
$ sudo mv cuda-wsl-ubuntu.pin /etc/apt/preferences.d/cuda-repository-pin-600
$ wget https://developer.download.nvidia.com/compute/cuda/11.8.0/local_installers/cuda-repo-wsl-ubuntu-11-8-local_11.8.0-1_amd64.deb
$ sudo dpkg -i cuda-repo-wsl-ubuntu-11-8-local_11.8.0-1_amd64.deb
$ sudo cp /var/cuda-repo-wsl-ubuntu-11-8-local/cuda-*-keyring.gpg /usr/share/keyrings/
$ sudo apt-get update
$ sudo apt-get -y install cuda
```

一定要选择WSL版本的

Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown. By downloading and using the software, you agree to fully comply with the terms and conditions of the [CUDA EULA](#).

Operating System

Linux

Windows

Architecture

x86_64

ppc64le

arm64-sbsa

aarch64-jetson

Distribution

CentOS

Debian

Fedora

KylinOS

OpenSUSE

RHEL

Rocky

SLES

Ubuntu

Version

WSL-Ubuntu

2.0

Installer Type

deb (local)

deb (network)

runfile (local)

Download Installer for Linux WSL-Ubuntu 2.0 x86_64

The base installer is available for download below.

Base Installer

Installation Instructions:

```
$ wget https://developer.download.nvidia.com/compute/cuda/repos/wsl-ubuntu/x86_64/cuda-wsl-ubuntu.pin
$ sudo mv cuda-wsl-ubuntu.pin /etc/apt/preferences.d/cuda-repository-pin-600
$ wget https://developer.download.nvidia.com/compute/cuda/11.8.0/local_installers/cuda-repo-wsl-ubuntu-11-8-local_11.8.0-1_amd64.deb
$ sudo dpkg -i cuda-repo-wsl-ubuntu-11-8-local_11.8.0-1_amd64.deb
$ sudo cp /var/cuda-repo-wsl-ubuntu-11-8-local/cuda-*-keyring.gpg /usr/share/keyrings/
$ sudo apt-get update
$ sudo apt-get -y install cuda
```

选择好后将nvidia官方给的命令，一条条复制到wsl2的terminal中即可。

安装结束之后执行nvcc -v，会提示没有nvcc可执行，这并不是因为我们cudatoolkit没安装好，而是因为环境变量还没配置好。

接下来我们将进行cuda环境变量配置：

1. 首先我们先在wsl2的terminal中输入

```
sudo nano ~/.bashrc
```

2. 将以下内容添加进文件最后

```
export PATH=/usr/local/cuda-11.8/bin${PATH:+:${PATH}}
export LD_LIBRARY_PATH=/usr/local/cuda-11.8/lib64${LD_LIBRARY_PATH:+:${LD_LIBRARY_PATH}}
```

3. 保存退出后（Ctrl+x再按 y 最后按回车），更新一下环境变量，输入：

```
source ~/.bashrc
```

4. 这时候在执行 nvcc -V 就能够显示cuda版本了。如图：

```
caicai@DESKTOP-26UFLT:~$ nvcc -V
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed Sep 21 10:33:58 PDT 2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda 11.8.r11.8/compiler.31833905_0
```

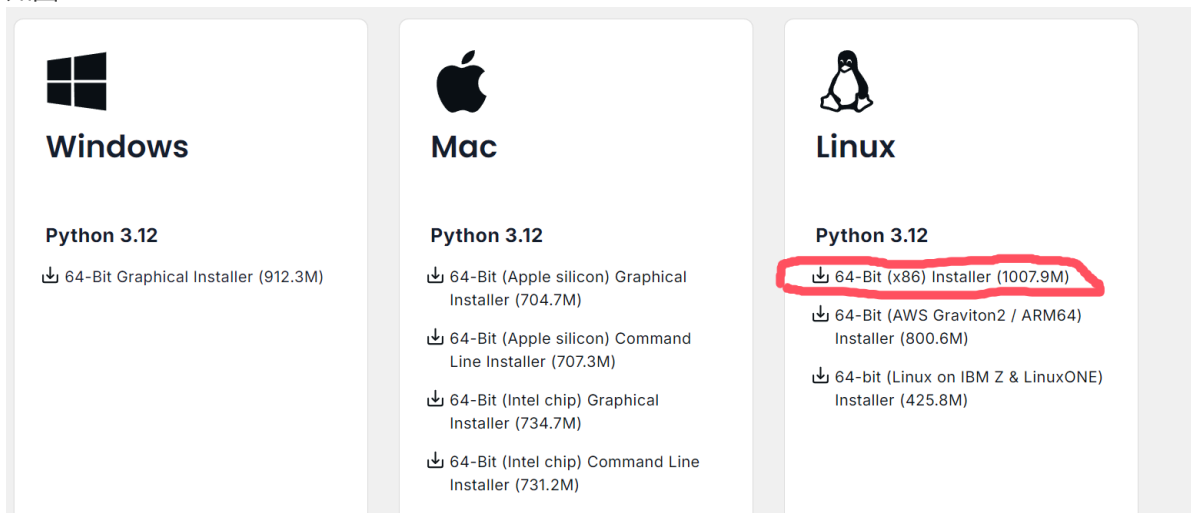
到此我们便完成了NVIDIA 驱动和 CUDA Toolkit的安装

wsl安装anaconda并配置环境以及PyTorch的安装

- 安装anaconda

[anaconda3官方下载](#),然后提交自己的邮箱,进入下载页面。选择linux版本,鼠标放在其上方右键,复制链接。

如图:



回到Ubuntu的terminal, 输入:

```
wget https://repo.anaconda.com/archive/Anaconda3-2024.06-1-Linux-x86_64.sh
```

我的版本是2024.06, 请根据自己当时复制的链接进行修改。

运行以上代码, 将会下载anaconda3到wsl ubuntu中。

之后按下sh A然后按Tab键, 系统会自动补齐下面内容。

如:

```
sh Anaconda3-2024.06-1-Linux-x86_64.sh
```

接下来就是安装过程, 只需要根据提示按回车或者输入yes即可。

- conda配置环境
用conda创建虚拟环境

```
conda create --name cu118py310 python=3.10 #--name 后面是创建环境的名字, 按自己的习惯命名, python=XX, 输入自己想用的版本号
```

例如我输入:

```
conda create --name cc python=3.10
```

当左侧出现(base)则说明成功了, 如:

```
(base) caicai@DESKTOP-26UFLTD:~$
```

当我们输入

```
conda info --env
```

我们则可以看到所有python环境，前面有个‘*’的代表当前环境

```
(base) caicai@DESKTOP-26UFLTD:~$ conda info --env
# conda environments:
#
base                *  /home/caicai/anaconda3
cc                  /home/caicai/anaconda3/envs/cc
```

比如我这个现在仍是base环境

我们需要通过

```
conda activate cc #（cc替换为我们创建的环境的名字） 激活刚刚创建的环境
```

然后我们可以明显看到我们切换到了cc环境，如图：

```
(cc) caicai@DESKTOP-26UFLTD:~$
```

- 配置pytorch

前往[pytorch官网](#)，选择需要的环境（注意这里选择linux OS），复制conda命令,在terminal中粘贴，回车，安装环境：

PyTorch Build	Stable (2.0.0)		Preview (Nightly)	
Your OS	Linux	Mac		Windows
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.7	CUDA 11.8	ROCm 5.4.2	CPU
Run this Command:	conda install pytorch torchvision torchaudio pytorch-cuda=11.8 -c pytorch -c nvidia			

我们还需要检验一下PyTorch是否安装成功，我们可以通过torch.cuda.is_available() 验证，首先输入python进入python环境再通过import torch导入torch库，回车后执行print(torch.cuda.is_available())再回车，如若输出True则代表我们成功了
以上操作如图所示：

```
(cc) caicai@DESKTOP-26UFLTD:~$ python
Python 3.10.15 (main, Oct 3 2024, 07:27:34) [GCC 11.2.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import torch
>>> print(torch.cuda.is_available())
True
>>> _
```

CIFAR-10 数据集的下载与处理

我们可以先与WSL的终端中使用:

```
code .
```

来进入vscode界面便于运行接下来的python代码

然后利用搜索引擎查找CIFAR-10 数据集的下载以及处理的代码，我们可以找到如下：

[参考](#)

```
import torchvision
import torchvision.transforms as transforms

# 数据预处理
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 加载数据集
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True,
transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True,
transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False)
```

通过以上便可实现该数据集的下载与处理，以及成功转变为 PyTorch DataLoader 格式，确保数据集可以高效加载到 GPU 进行训练。

实现深度学习模型

根据考核要求：使用 PyTorch 实现一个基础的卷积神经网络（CNN），模型结构可以包括卷积层、池化层和全连接层，不调用现成的模型库（如 torchvision.models）。并在 GPU 上训练该模型，并验证其性能。

为了验证模型的性能，我们可以以以下标准进行测量：

1. 准确率（Accuracy）：
这是衡量模型性能的最基本指标，表示模型正确预测的样本数量占总样本数量的比例。准确率越高，说明模型的预测能力越强。
2. 混淆矩阵（Confusion Matrix）：
混淆矩阵是一个表格，用于显示模型预测的结果与实际标签之间的关系。它可以帮助我们了解模型在各个类别上的表现，特别是哪些类别容易被混淆。
3. 精确率（Precision）：精确率是指模型预测为正类别中实际为正类别的比例。对于每个类别，精确率告诉我们模型预测的准确性。
4. 召回率（Recall）：召回率是指在所有实际为正类别的样本中，模型正确预测为正类别的比例。召回率越高，说明模型漏掉的正样本越少。
5. F1分数（F1 Score）：F1分数是精确率和召回率的调和平均数，它综合考虑了精确率和召回率的平衡。F1分数越高，说明模型在精确率和召回率之间取得了较好的平衡。
6. 处理所有测试数据所需的时间（Training Time）：用来检测模型的运行快慢
我们不妨通过以上标准来评估模型的性能，包括模型的整体效果、在不同类别上的表现、以及模型的稳定性和可靠性。通过这些指标，我们可以对模型的优缺点有一个清晰的认识，并据此进行模型的优化和调整。
7. 平均内存使用量（Average Memory Usage）：反应模型的内存效率

我们的Python环境中没有安装scikit-learn库。这个库提供了许多用于评估模型性能的工具。我们可以通过以下命令安装：

```
conda install scikit-learn
```

然后便可以继续运行，
以下直接给出注有详细注释的代码：

```
import os
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import time
import psutil
from sklearn.metrics import confusion_matrix, precision_recall_fscore_support

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=64, shuffle=True, num_workers=4)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=64, shuffle=False, num_workers=4)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        # 第一个卷积层，输入通道3（RGB图像），输出通道32，卷积核大小3x3
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        # 池化层，窗口大小2x2，步长2
        self.pool = nn.MaxPool2d(2, 2)
        # 第二个卷积层，输入通道32，输出通道64，卷积核大小3x3
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        # 第一个全连接层，输入特征数为64*8*8（因为经过两次池化后，特征图大小减半两次），输出特
        # 征数256
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        # 第二个全连接层，输入特征数256，输出特征数10（CIFAR-10数据集的类别数）
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        # 应用第一个卷积层和激活函数ReLU
```



```

x = self.pool(F.relu(self.conv1(x)))
# 应用第二个卷积层和激活函数ReLU
x = self.pool(F.relu(self.conv2(x)))
# 展平特征图，为全连接层准备
x = x.view(-1, 64 * 8 * 8)
# 应用第一个全连接层和激活函数ReLU
x = F.relu(self.fc1(x))
# 应用第二个全连接层
x = self.fc2(x)
return x

def train_and_evaluate(device):
    net = Net().to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
    criterion = nn.CrossEntropyLoss().to(device)

    start_time = time.time()

    for epoch in range(10): # 使用10个epoch作为示例
        running_loss = 0.0
        for i, (inputs, labels) in enumerate(trainloader, 0):
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()

            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            if i % 2000 == 1999: # 每2000个小批量打印一次
                print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 2000:.3f}')
                running_loss = 0.0
                print(f'Memory Usage:
{psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f} MB') # 打印内存使用情况

        print('Finished Training')

    end_time = time.time()
    elapsed_time = end_time - start_time
    print(f'Training took {elapsed_time:.2f} seconds on {device}')

    # 测试模型性能
    all_labels = []
    all_preds = []
    all_probs = []
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in testloader:
            images = images.to(device)
            labels = labels.to(device)
            outputs = net(images)
            _, predicted = torch.max(outputs.data, 1)

```

```

        total += labels.size(0)
        correct += (predicted == labels).sum().item()
        all_labels.extend(labels.cpu().numpy())
        all_preds.extend(predicted.cpu().numpy())
        all_probs.extend(outputs.cpu().numpy())

    accuracy = 100 * correct / total
    cm = confusion_matrix(all_labels, all_preds)
    precision, recall, f1, _ = precision_recall_fscore_support(all_labels,
all_preds, average='weighted')

    return elapsed_time, accuracy, cm, precision, recall, f1,
psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024)

# 训练和评估在CPU上
cpu_time, cpu_accuracy, cpu_cm, cpu_precision, cpu_recall, cpu_f1,
cpu_memory_usage = train_and_evaluate(torch.device("cpu"))
print(f'CPU - Training Time: {cpu_time:.2f} seconds, Accuracy:
{cpu_accuracy:.2f}%, Precision: {cpu_precision:.2f}, Recall: {cpu_recall:.2f}, F1
Score: {cpu_f1:.2f}, Memory Usage: {cpu_memory_usage:.2f} MB')
print(f'Confusion Matrix: \n{cpu_cm}')

# 训练和评估在GPU上
if torch.cuda.is_available():
    gpu_time, gpu_accuracy, gpu_cm, gpu_precision, gpu_recall, gpu_f1,
gpu_memory_usage = train_and_evaluate(torch.device("cuda"))
    print(f'GPU - Training Time: {gpu_time:.2f} seconds, Accuracy:
{gpu_accuracy:.2f}%, Precision: {gpu_precision:.2f}, Recall: {gpu_recall:.2f}, F1
Score: {gpu_f1:.2f}, Memory Usage: {gpu_memory_usage:.2f} MB')
    print(f'Confusion Matrix: \n{gpu_cm}')
else:
    print("CUDA is not available. Cannot compare with GPU.")

```

我们尝试将模型分别在CPU和GPU加速下进行比较，并且设置了batch size=64，workers=4。运行，我们可以得到我们想要的结果，如我所运行的一次结果：

```

Files already downloaded and verified
Files already downloaded and verified
Training set size: 50000
Test set size: 10000
Finished Training
Training took 83.57 seconds on cpu
CPU - Training Time: 83.57 seconds, Accuracy: 62.94%, Precision: 0.64, Recall: 0.63, F1 Score: 0.62, Memory Usage: 783.43 MB
Confusion Matrix:
[[780 12 34 10 14 4 14 3 110 19]
 [ 55 701 12 4 7 2 10 3 95 111]
 [117 6 430 47 203 54 62 34 34 13]
 [ 52 6 81 377 164 108 101 40 50 21]
 [ 48 2 64 42 709 16 47 44 26 2]
 [ 31 3 92 160 107 456 45 67 30 9]
 [ 23 4 36 44 132 11 713 8 20 9]
 [ 44 0 28 39 134 45 17 658 8 27]
 [ 93 24 10 7 7 3 6 2 831 17]
 [ 76 113 11 12 15 4 11 19 100 639]]
Finished Training
Training took 21.83 seconds on cuda
GPU - Training Time: 21.83 seconds, Accuracy: 63.35%, Precision: 0.64, Recall: 0.63, F1 Score: 0.63, Memory Usage: 4660.57 MB
Confusion Matrix:
[[805 17 27 10 14 7 8 6 50 56]
 [ 42 703 10 7 3 3 5 2 32 193]
 [116 11 446 62 172 70 41 37 23 22]
 [ 51 9 61 386 119 200 65 34 30 45]
 [ 55 2 71 40 661 41 43 54 17 16]
 [ 30 3 71 147 88 542 31 53 16 19]
 [ 21 10 52 60 110 26 666 11 10 34]
 [ 34 2 24 40 118 70 10 649 3 50]
 [145 50 9 12 7 4 1 4 719 49]
 [ 57 98 6 8 9 8 9 8 39 758]]

```

在GPU加速下，速度明显更快，但两者的精度都维持在63%上下。

使用 GPU 加速，优化模型训练速度

• 调整 batch size

我们定义了一个列表batch_sizes，它包含了我们想要测试的不同批量大小。对于每个批量大小，我们创建了新的DataLoader实例，并记录了训练模型所需的总时间。然后，我们计算了训练速度（以样本/秒为单位），并在测试集上评估了模型的精度。给出代码：

```
import os
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import time
import psutil

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

def train_and_evaluate(device, batch_size):
    net = Net().to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```

criterion = nn.CrossEntropyLoss().to(device)
trainloader = DataLoader(trainset, batch_size=batch_size, shuffle=True,
num_workers=2)
testloader = DataLoader(testset, batch_size=batch_size, shuffle=False,
num_workers=2)

start_time = time.time()

for epoch in range(10): # 使用10个epoch作为示例
    running_loss = 0.0
    for i, (inputs, labels) in enumerate(trainloader, 0):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        if i % 100 == 99: # 每100个小批量打印一次
            print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 100:.3f}')
            running_loss = 0.0

    print('Finished Training')

end_time = time.time()
elapsed_time = end_time - start_time
print(f'Training Time (seconds) for batch_size {batch_size}:
{elapsed_time:.2f}')

# 测试模型性能
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0
print(f'Accuracy (%) for batch_size {batch_size}: {accuracy:.2f}')

memory_usage = psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024)
print(f'Memory Usage (MB) for batch_size {batch_size}: {memory_usage:.2f}')

return elapsed_time, accuracy, memory_usage

# 训练和评估在GPU上
if torch.cuda.is_available():
    batch_sizes = [16, 32, 64, 128] # 不同的batch_size进行比较
    for batch_size in batch_sizes:

```

```
train_time, accuracy, memory_usage =  
train_and_evaluate(torch.device("cuda"), batch_size)  
else:  
    print("CUDA is not available. Cannot train on GPU.")
```

运行，得到结果：

1. 当batch_size=16时：

```
Training Time (seconds) for batch_size 16: 63.72  
Accuracy (%) for batch_size 16: 71.73  
Memory Usage (MB) for batch_size 16: 4576.65
```

2. 当batch_size=32时：

```
Training Time (seconds) for batch_size 32: 33.94  
Accuracy (%) for batch_size 32: 69.47  
Memory Usage (MB) for batch_size 32: 4579.46
```

3. 当batch_size=64时：

```
Training Time (seconds) for batch_size 64: 25.62  
Accuracy (%) for batch_size 64: 62.45  
Memory Usage (MB) for batch_size 64: 4579.64
```

4. 当batch_size=128时：

```
Training Time (seconds) for batch_size 128: 23.20  
Accuracy (%) for batch_size 128: 56.87  
Memory Usage (MB) for batch_size 128: 4580.03
```

不难看出，随着batch_size的不断增高训练用时越来越短，精度逐渐降低，而内存的占用几乎没有明显的变化

• 使用混合精度训练 (torch.cuda.amp)

先给出不使用该优化的代码：

```
import os  
import torch  
import torchvision  
import torchvision.transforms as transforms  
from torch.utils.data import DataLoader  
import torch.nn as nn  
import torch.optim as optim  
import torch.nn.functional as F  
import time  
import psutil  
  
# 数据集放置路径  
data_save_path = "./data"  
  
# 数据转换  
transform = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))  
)  
  
# 使用官方方式加载数据集
```

```

trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

def train_and_evaluate(device):
    net = Net().to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
    criterion = nn.CrossEntropyLoss().to(device)

    start_time = time.time()

    for epoch in range(10): # 使用10个epoch作为示例
        running_loss = 0.0
        for i, (inputs, labels) in enumerate(trainloader, 0):
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()

            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            if i % 2000 == 1999: # 每2000个小批量打印一次
                print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 2000:.3f}')
                running_loss = 0.0
                print(f'Memory Usage:
{psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f} MB')

```

```

print('Finished Training')

end_time = time.time()
elapsed_time = end_time - start_time

# 测试模型性能
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0

print(f'Training Time (seconds): {elapsed_time:.2f}')
print(f'Accuracy (%): {accuracy:.2f}')
print(f'Memory Usage (MB): {psutil.Process(os.getpid()).memory_info().rss /
(1024 * 1024):.2f}')

return elapsed_time, accuracy, psutil.Process(os.getpid()).memory_info().rss /
(1024 * 1024)

# 训练和评估在GPU上
if torch.cuda.is_available():
    train_and_evaluate(torch.device("cuda"))
else:
    print("CUDA is not available. Cannot train on GPU.")

```

结果如下:

```

Finished Training
Accuracy of the network on the test images: 72.94%
Training Time (seconds): 230.00
Initial Memory Usage (MB): 2434.84
Final Memory Usage (MB): 4589.73

```

使用了混合精度训练 (torch.cuda.amp) 时, 代码:

```

import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import time
import psutil
import os

# 数据集放置路径
data_save_pth = "./data"

# 数据转换

```

```

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return x

# 实例化模型、优化器、损失函数和GradScaler
model = Net().cuda()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
criterion = nn.CrossEntropyLoss().cuda()
scaler = torch.amp.GradScaler() # 使用新的API

# 训练和评估模型
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
model.train()

start_time = time.time()
memory_usage_start = psutil.Process(os.getpid()).memory_info().rss / (1024 *
1024) # 记录初始内存使用

for epoch in range(10): # 使用10个epoch作为示例
    for i, (inputs, labels) in enumerate(trainloader, 0):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

```



```

# 使用autocast进行混合精度计算
with torch.amp.autocast(device_type='cuda'):
    outputs = model(inputs)
    loss = criterion(outputs, labels)

# 使用GradScaler来缩放梯度
scaler.scale(loss).backward()
scaler.step(optimizer)
scaler.update()

if i % 2000 == 1999: # 每2000个小批量打印一次
    print(f'[{epoch + 1}, {i + 1}] loss: {loss.item():.3f}')
    print(f'Memory Usage: {psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f} MB')

end_time = time.time()
elapsed_time = end_time - start_time
print('Finished Training')

# 测试模型性能
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0
print(f'Accuracy of the network on the test images: {accuracy:.2f}%')
print(f'Training Time (seconds): {elapsed_time:.2f}')
print(f'Initial Memory Usage (MB): {memory_usage_start:.2f}')
print(f'Final Memory Usage (MB): {psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f}')

```

```

Accuracy of the network on the test images: 71.56%
Training Time (seconds): 276.97
Initial Memory Usage (MB): 730.50
Final Memory Usage (MB): 1120.71

```

比较可得，使用混合精度训练使得用时和内存占用显著降低，但是精度反而下降了点。经过调查得到原因如下：

1. 数值稳定性问题：FP16的数值范围比FP32小，可能会导致数值下溢（underflow）或精度损失。这可能会在训练过程中引入舍入误差，影响模型的收敛和最终性能。
2. 梯度溢出或下溢：在FP16下，梯度的数值可能变得非常小，以至于在FP16格式下无法有效表示，这可能导致梯度被置为零，从而影响模型的学习。

3. 模型对精度敏感：某些模型或任务对数值精度非常敏感。在这些情况下，混合精度训练可能会因为精度损失而导致模型性能下降。
 - 。
4. 硬件兼容性：不是所有的硬件都支持FP16运算。在没有专门Tensor Core的GPU上，使用FP16可能不会带来预期的性能提升，甚至可能影响模型的精度。
 - 。
5. 调试和分析困难：使用混合精度训练可能会使得模型的调试和性能分析更加复杂，因为需要跟踪哪些操作是在FP16下执行的，哪些是在FP32下执行的。这种复杂性可能导致错误的配置或实现，从而影响模型的性能。
 - 。
6. 模型泛化能力：在某些情况下，混合精度训练可能会影响模型的泛化能力，尤其是在模型对精度非常敏感的情况下。因此，可能需要对模型进行微调以确保其性能和精度的稳定性。
 - 。

为了优化混合精度训练并减少精度损失，我们可以采取一些措施，比如使用 `torch.amp.autocast` 和 `torch.amp.GradScaler` 来实现自动混合精度训练，`autocast` 会自动选择使用 FP16 或 FP32 进行计算，以提高性能并减少内存使用，通过 `GradScaler` 的动态缩放，可以减少由于 FP16 精度限制导致的数值稳定性问题。代码如下：

```
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
from torch.amp import autocast, GradScaler # 正确的导入方式
import time
import psutil
import os

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
```

```

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return x

# 实例化模型、优化器、损失函数和GradScaler
model = Net().cuda()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
criterion = nn.CrossEntropyLoss().cuda()
scaler = GradScaler() # 初始化GradScaler

# 训练和评估模型
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
model.train()

start_time = time.time()
memory_usage_start = psutil.Process(os.getpid()).memory_info().rss / (1024 *
1024) # 记录初始内存使用

for epoch in range(10): # 使用10个epoch作为示例
    for i, (inputs, labels) in enumerate(trainloader, 0):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        # 使用autocast进行混合精度计算
        with autocast(device_type='cuda'): # 正确的autocast使用方式
            outputs = model(inputs)
            loss = criterion(outputs, labels)

        # 使用GradScaler来缩放梯度
        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

        if i % 2000 == 1999: # 每2000个小批量打印一次
            print(f'[{epoch + 1}, {i + 1}] loss: {loss.item():.3f}')
            print(f'Memory Usage: {psutil.Process(os.getpid()).memory_info().rss
/ (1024 * 1024):.2f} MB')

end_time = time.time()
elapsed_time = end_time - start_time
print('Finished Training')

```

```

# 测试模型性能
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0
print(f'Accuracy of the network on the test images: {accuracy:.2f}%')
print(f'Training Time (seconds): {elapsed_time:.2f}')
print(f'Initial Memory Usage (MB): {memory_usage_start:.2f}')
print(f'Final Memory Usage (MB): {psutil.Process(os.getpid()).memory_info().rss /
(1024 * 1024):.2f}')

```

• 使用混合精度训练 (torch.cuda.amp) 下, 调整 batch size

结合以上两者, 我们不妨直接给出测试代码:

```

import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
from torch.amp import autocast, GradScaler # 正确的导入方式
import time
import psutil
import os

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)

```

```

self.pool = nn.MaxPool2d(2, 2)
self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
self.fc1 = nn.Linear(64 * 8 * 8, 256)
self.fc2 = nn.Linear(256, 10)

def forward(self, x):
    x = self.pool(torch.relu(self.conv1(x)))
    x = self.pool(torch.relu(self.conv2(x)))
    x = x.view(-1, 64 * 8 * 8)
    x = torch.relu(self.fc1(x))
    x = self.fc2(x)
    return x

# 实例化模型、优化器和损失函数
model = Net().cuda()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
criterion = nn.CrossEntropyLoss().cuda()
scaler = GradScaler() # 初始化GradScaler

# 训练和评估模型
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)

# 测试不同的batch size
batch_sizes = [32, 64, 128, 256] # 不同的batch size
for batch_size in batch_sizes:
    trainloader = DataLoader(trainset, batch_size=batch_size, shuffle=True,
num_workers=2)
    testloader = DataLoader(testset, batch_size=batch_size, shuffle=False,
num_workers=2)

    model.train()
    start_time = time.time()
    memory_usage_start = psutil.Process(os.getpid()).memory_info().rss / (1024 *
1024)

    for epoch in range(10): # 使用10个epoch作为示例
        for i, (inputs, labels) in enumerate(trainloader, 0):
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()

            with autocast(device_type='cuda'): # 正确的autocast使用方式
                outputs = model(inputs)
                loss = criterion(outputs, labels)

            scaler.scale(loss).backward()
            scaler.step(optimizer)
            scaler.update()

    end_time = time.time()
    elapsed_time = end_time - start_time
    print(f'Batch size: {batch_size}')
    print(f'Training Time (seconds): {elapsed_time:.2f}')
    print(f'Initial Memory Usage (MB): {memory_usage_start:.2f}')
    print(f'Final Memory Usage (MB):
{psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f}')

```

```

# 测试模型性能
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0
print(f'Accuracy of the network on the test images: {accuracy:.2f}%\n')

```

```

Batch size: 32
Training Time (seconds): 42.74
Initial Memory Usage (MB): 732.80
Final Memory Usage (MB): 1092.00
Accuracy of the network on the test images: 69.92%

Batch size: 64
Training Time (seconds): 26.61
Initial Memory Usage (MB): 1130.07
Final Memory Usage (MB): 1134.13
Accuracy of the network on the test images: 71.38%

Batch size: 128
Training Time (seconds): 24.12
Initial Memory Usage (MB): 1135.01
Final Memory Usage (MB): 1147.59
Accuracy of the network on the test images: 71.82%

Batch size: 256
Training Time (seconds): 27.74
Initial Memory Usage (MB): 1147.66
Final Memory Usage (MB): 1147.73
Accuracy of the network on the test images: 72.25%

```

得到结果如下：

可以看出：较大的 batch size 可以提高 GPU 的利用率，因为它允许更高效地进行并行计算；更大的 batch size 意味着每个 epoch 中的迭代次数减少，从而减少了训练所需的总迭代次数；同时较大的 batch size 会占用更多的显存。

• 使用学习率调度器 (Learning Rate Scheduler)

为了比较，我们先给出不使用学习率调度器的代码，并运行：

```

import os
import torch
import torchvision

```

```

import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import time
import psutil

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

def train_and_evaluate(device):
    net = Net().to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
    criterion = nn.CrossEntropyLoss().to(device)

    start_time = time.time()

    for epoch in range(10): # 使用10个epoch作为示例

```

```

running_loss = 0.0
for i, (inputs, labels) in enumerate(trainloader, 0):
    inputs, labels = inputs.to(device), labels.to(device)

    optimizer.zero_grad()

    outputs = net(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    running_loss += loss.item()
    if i % 2000 == 1999: # 每2000个小批量打印一次
        print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 2000:.3f}')
        running_loss = 0.0
        print(f'Memory Usage:
{psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f} MB')

    print('Finished Training')

end_time = time.time()
elapsed_time = end_time - start_time

# 测试模型性能
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0

print(f'Training Time (seconds): {elapsed_time:.2f}')
print(f'Accuracy (%): {accuracy:.2f}')
print(f'Memory Usage (MB): {psutil.Process(os.getpid()).memory_info().rss /
(1024 * 1024):.2f}')

return elapsed_time, accuracy, psutil.Process(os.getpid()).memory_info().rss
/ (1024 * 1024)

# 训练和评估在GPU上
if torch.cuda.is_available():
    train_and_evaluate(torch.device("cuda"))
else:
    print("CUDA is not available. Cannot train on GPU.")

```

得到结果如下:

```

Training Time (seconds): 283.13
Accuracy (%): 71.76
Memory Usage (MB): 4596.39

```

接下来我们再使用学习率调度器，直接贴出代码：


```

import os
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import time
import psutil

# 数据集放置路径
data_save_pth = "./data"

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 使用官方方式加载数据集
trainset = torchvision.datasets.CIFAR10(root=data_save_pth, train=True,
download=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root=data_save_pth, train=False,
download=True, transform=transform)
testloader = DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

# 检查数据集大小
print(f'Training set size: {len(trainset)}')
print(f'Test set size: {len(testset)}')

# 定义CNN模型
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, 3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 64, 3, stride=1, padding=1)
        self.fc1 = nn.Linear(64 * 8 * 8, 256)
        self.fc2 = nn.Linear(256, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

def train_and_evaluate(device):
    net = Net().to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
    criterion = nn.CrossEntropyLoss().to(device)

```

```

scaler = torch.amp.GradScaler() # 初始化GradScaler

start_time = time.time()

for epoch in range(10): # 使用10个epoch作为示例
    running_loss = 0.0
    for i, (inputs, labels) in enumerate(trainloader, 0):
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        with torch.amp.autocast(device_type="cuda"): # 开启自动混合精度
            outputs = net(inputs)
            loss = criterion(outputs, labels)

        scaler.scale(loss).backward() # 缩放损失以避免梯度下溢
        scaler.step(optimizer) # 更新参数
        scaler.update() # 更新缩放器

        running_loss += loss.item()
        if i % 2000 == 1999: # 每2000个小批量打印一次
            print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 2000:.3f}')
            running_loss = 0.0
            print(f'Memory Usage:
{psutil.Process(os.getpid()).memory_info().rss / (1024 * 1024):.2f} MB')

    print('Finished Training')

end_time = time.time()
elapsed_time = end_time - start_time

# 测试模型性能
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total if total > 0 else 0

print(f'Training Time (seconds): {elapsed_time:.2f}')
print(f'Accuracy (%): {accuracy:.2f}')
print(f'Memory Usage (MB): {psutil.Process(os.getpid()).memory_info().rss /
(1024 * 1024):.2f}')

return elapsed_time, accuracy, psutil.Process(os.getpid()).memory_info().rss
/ (1024 * 1024)

# 训练和评估在GPU上
if torch.cuda.is_available():
    train_and_evaluate(torch.device("cuda"))

```

```
else:
    print("CUDA is not available. Cannot train on GPU.")
```

运行得到结果如下：

```
Training Time (seconds): 221.61
Accuracy (%): 75.25
Memory Usage (MB): 4579.92
```

不难看出，虽然用时变短，内存占用几乎没变化，但精确度确实存在一定的提高。

使用cuDNN

安装cuDNN，我们可以进入[官网](#)

按照所给命令在终端上输入即可，如我的：

```
wget
https://developer.download.nvidia.com/compute/cudnn/9.5.0/local_installers/cudnn-
local-repo-ubuntu2204-9.5.0_1.0-1_amd64.deb

sudo dpkg -i cudnn-local-repo-ubuntu2204-9.5.0_1.0-1_amd64.deb

sudo cp /var/cudnn-local-repo-ubuntu2204-9.5.0/cudnn-*-keyring.gpg
/usr/share/keyrings/

sudo apt-get update

sudo apt-get -y install cudnn
```

通过以上指令，我们成功安装了cuDNN，接下来我们可以使用cuDNN来训练模型。

以下为代码：

```
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
import time

# 定义CNN模型
class SimpleCNN(nn.Module):
    def __init__(self):
        super(SimpleCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(64 * 8 * 8, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
```

```

        return x

# 数据集路径
data_dir = './data'

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 加载CIFAR-10数据集
train_dataset = torchvision.datasets.CIFAR10(root=data_dir, train=True,
                                              download=True, transform=transform)
test_dataset = torchvision.datasets.CIFAR10(root=data_dir, train=False,
                                              download=True, transform=transform)

# 创建数据加载器
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True,
                           num_workers=2)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False,
                          num_workers=2)

# 设置设备
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# 初始化模型、优化器和损失函数
model = SimpleCNN().to(device)
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

# 启用cuDNN
torch.backends.cudnn.enabled = True
torch.backends.cudnn.benchmark = True

# 训练模型
num_epochs = 10
for epoch in range(num_epochs):
    model.train()
    start_time = time.time()
    total_loss = 0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    end_time = time.time()
    print(f'Epoch {epoch+1}, Loss: {total_loss / len(train_loader):.4f}, Time: {end_time - start_time:.4f} seconds')

# 测试模型性能
model.eval()
correct = 0

```

```
total = 0
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = correct / total
print(f'Test Accuracy: {accuracy:.4f}')
```

运行得到以下结果：

```
Files already downloaded and verified
Files already downloaded and verified
Epoch 1, Loss: 1.3454, Time: 3.5468 seconds
Epoch 2, Loss: 0.9649, Time: 2.4095 seconds
Epoch 3, Loss: 0.8063, Time: 2.5018 seconds
Epoch 4, Loss: 0.6838, Time: 2.5234 seconds
Epoch 5, Loss: 0.5761, Time: 2.4362 seconds
Epoch 6, Loss: 0.4823, Time: 2.4228 seconds
Epoch 7, Loss: 0.3945, Time: 2.4565 seconds
Epoch 8, Loss: 0.3136, Time: 2.4530 seconds
Epoch 9, Loss: 0.2399, Time: 2.5430 seconds
Epoch 10, Loss: 0.1885, Time: 2.6930 seconds
Test Accuracy: 0.7165
```

性能提升并不明显，可能原因如下：

1. GPU 没有被充分利用，可能是因为批量大小太小，或者模型的并行化程度不够。
2. 模型相对较小，或者数据集不够大，那么 GPU 和 cuDNN 提供的加速效果可能不会非常明显。
3. 系统上同时运行了其他资源密集型的进程，它们可能会与你的训练任务竞争 CPU、内存或 I/O 资源，从而影响性能。

但当我们调整batch size和workers时可以发现，速度有了明显上升。所以，cuDNN的加速在速度方面有显著作用。

不同优化方法分别在CPU和GPU上测试的性能比较

• 调整 batch size下的比较

```
import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.optim as optim
from torch.amp import autocast, GradScaler
import time

# 定义CNN模型
class SimpleCNN(nn.Module):
    def __init__(self):
```

```

        super(SimpleCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(64 * 8 * 8, 512)
        self.fc2 = nn.Linear(512, 10)

    def forward(self, x):
        x = self.pool(torch.relu(self.conv1(x)))
        x = self.pool(torch.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return x

# 数据集路径
data_dir = './data'

# 数据转换
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# 加载CIFAR-10数据集
train_dataset = torchvision.datasets.CIFAR10(root=data_dir, train=True,
                                              download=True, transform=transform)
test_dataset = torchvision.datasets.CIFAR10(root=data_dir, train=False,
                                              download=True, transform=transform)

# 创建数据加载器
def create_loader(dataset, batch_size, shuffle=True):
    return DataLoader(dataset, batch_size=batch_size, shuffle=shuffle,
                      num_workers=2)

# 初始化模型、优化器和损失函数
def init_model(device, batch_size):
    model = SimpleCNN().to(device)
    optimizer = optim.Adam(model.parameters(), lr=0.001)
    criterion = nn.CrossEntropyLoss()
    scaler = GradScaler()
    train_loader = create_loader(train_dataset, batch_size)
    test_loader = create_loader(test_dataset, batch_size, shuffle=False)
    return model, optimizer, criterion, scaler, train_loader, test_loader

# 训练模型
def train_model(model, device, train_loader, optimizer, criterion, scaler):
    model.train()
    total_loss = 0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer.zero_grad()

        # 使用自动混合精度, 指定device_type
        if device.type == 'cuda':
            with autocast(device_type='cuda'):

```

```

        outputs = model(inputs)
        loss = criterion(outputs, labels)
    else:
        outputs = model(inputs)
        loss = criterion(outputs, labels)

    scaler.scale(loss).backward()
    scaler.step(optimizer)
    scaler.update()
    total_loss += loss.item()
    return total_loss / len(train_loader)

# 测试模型性能
def test_model(model, device, test_loader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for inputs, labels in test_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    return correct / total

# 主函数
def main():
    batch_sizes = [32, 64, 128] # 测试不同的批量大小
    for batch_size in batch_sizes:
        print(f"Testing with batch size: {batch_size}")

        # 在CPU上测试
        device = torch.device("cpu")
        model_cpu, optimizer_cpu, criterion_cpu, scaler_cpu, train_loader_cpu,
        test_loader_cpu = init_model(device, batch_size)
        start_time_cpu = time.time()
        train_loss_cpu = train_model(model_cpu, device, train_loader_cpu,
        optimizer_cpu, criterion_cpu, scaler_cpu)
        test_acc_cpu = test_model(model_cpu, device, test_loader_cpu)
        print(f"CPU - Loss: {train_loss_cpu:.4f}, Accuracy: {test_acc_cpu:.4f},
        Time: {time.time() - start_time_cpu:.4f} seconds")

        # 在GPU上测试
        if torch.cuda.is_available():
            device = torch.device("cuda")
            model_gpu, optimizer_gpu, criterion_gpu, scaler_gpu,
            train_loader_gpu, test_loader_gpu = init_model(device, batch_size)
            start_time_gpu = time.time()
            train_loss_gpu = train_model(model_gpu, device, train_loader_gpu,
            optimizer_gpu, criterion_gpu, scaler_gpu)
            test_acc_gpu = test_model(model_gpu, device, test_loader_gpu)
            print(f"GPU - Loss: {train_loss_gpu:.4f}, Accuracy:
            {test_acc_gpu:.4f}, Time: {time.time() - start_time_gpu:.4f} seconds")

if __name__ == "__main__":

```

```
main()
```

运行得到结果：

```
Files already downloaded and verified
Files already downloaded and verified
Testing with batch size: 32
CPU - Loss: 1.2394, Accuracy: 0.6178, Time: 33.4057 seconds
GPU - Loss: 1.2152, Accuracy: 0.6636, Time: 10.3509 seconds
Testing with batch size: 64
CPU - Loss: 1.2986, Accuracy: 0.6335, Time: 23.8536 seconds
GPU - Loss: 1.2635, Accuracy: 0.6398, Time: 4.9362 seconds
Testing with batch size: 128
CPU - Loss: 1.3606, Accuracy: 0.6063, Time: 18.3343 seconds
GPU - Loss: 1.3639, Accuracy: 0.6132, Time: 4.4172 seconds
```

可以看出无论batch size多少，在GPU的加速下，精度和速度都优于CPU。

通过调查所知道的原因：

1. 并行处理能力：GPU设计用于同时处理大量数据，这使得它们在执行并行计算时非常高效。GPU拥有成百上千个核心，可以同时处理多个计算任务，而CPU通常只有几个核心。
2. 浮点运算性能：GPU在浮点运算方面通常比CPU更快，这对于深度学习、科学计算和图形渲染等任务至关重要。
3. 内存带宽：GPU通常具有更高的内存带宽，这意味着它们可以更快地读取和写入大量数据，这对于处理大型数据集和复杂模型非常有用。

.....

等。

使实验数据可视化

先将算出的数据放入csv文件中，如下：

```
Optimization Technique,Training Time (seconds),Accuracy (%),Memory Usage (MB)
C_b=32_w=4,98.02,67.59,698.39
G_b=32_w=4,35.66,68.02,4600.46
C_b=64_w=4,84.21,61.31,735.86
G_b=64_w=4,22.13,61.25,4607.96
C_b=128_w=4,84.51,53.02,752.80
G_b=128_w=4,17.03,53.39,4622.02
C_b=128_w=6,81.49,53.38,797.48
G_b=128_w=6,17.56,53.45,4612.20
C_b=64_w=6,85.63,61.39,742.11
G_b=64_w=6,22.11,61.89,4603.28
C_torch.cuda.amp,479.67,72.87,667.46
G_torch.cuda.amp,359.58,71.90,4559.08
G_cuDNN,234.44,71.34,4586.90
```

其中G代表GPU，C代表CPU，b代表batch_size,w代表workers。

接下来我们使用python环境使得数据可视化，我们将会用到Python环境中的pandas库，matplotlib库和seaborn库。如若未安装可使用以下命令来安装：


```
pip install pandas -i https://pypi.tuna.tsinghua.edu.cn/simple
pip install matplotlib -i https://pypi.tuna.tsinghua.edu.cn/simple
pip install -i https://pypi.tuna.tsinghua.edu.cn/simple seaborn
```

接下来通过代码来让电脑进行绘制：

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# 设置Seaborn的样式
sns.set(style="whitegrid")

# 读取CSV文件
df = pd.read_csv('data.csv')

# 设置图表大小
plt.figure(figsize=(12, 10))

# 定义颜色列表
colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd', '#8c564b']

# 绘制Training Time柱状图
plt.subplot(3, 1, 1)
sns.barplot(x='Optimization Technique', y='Training Time (seconds)', data=df,
            palette=colors)
plt.title('Training Time by Optimization Technique', fontsize=16)
plt.xlabel('Optimization Technique', fontsize=14)
plt.ylabel('Training Time (seconds)', fontsize=14)
plt.xticks(rotation=45) # 设置横轴标签旋转45度

# 在柱状图上添加数据标签
for p in plt.gca().patches:
    plt.text(p.get_x() + p.get_width() / 2., p.get_height(),
            f'{p.get_height():.2f}',
            ha='center', va='bottom')

# 绘制Accuracy柱状图
plt.subplot(3, 1, 2)
sns.barplot(x='Optimization Technique', y='Accuracy (%)', data=df,
            palette=colors)
plt.title('Accuracy by Optimization Technique', fontsize=16)
plt.xlabel('Optimization Technique', fontsize=14)
plt.ylabel('Accuracy (%)', fontsize=14)
plt.xticks(rotation=45) # 设置横轴标签旋转45度

# 在柱状图上添加数据标签
for p in plt.gca().patches:
    plt.text(p.get_x() + p.get_width() / 2., p.get_height(),
            f'{p.get_height()}%',
            ha='center', va='bottom')

# 绘制Memory Usage柱状图
plt.subplot(3, 1, 3)
sns.barplot(x='Optimization Technique', y='Memory Usage (MB)', data=df,
```

```

        palette=colors)
plt.title('Memory Usage by Optimization Technique', fontsize=16)
plt.xlabel('Optimization Technique', fontsize=14)
plt.ylabel('Memory Usage (MB)', fontsize=14)
plt.xticks(rotation=45) # 设置横轴标签旋转45度

# 在柱状图上添加数据标签
for p in plt.gca().patches:
    plt.text(p.get_x() + p.get_width() / 2., p.get_height(),
             f'{p.get_height():.2f} MB',
             ha='center', va='bottom')


# 调整子图间距
plt.tight_layout()

# 保存图表
plt.savefig('optimization_techniques_comparison.png', dpi=300)

# 显示图表
plt.show()

```

得到结果如下：



实验中遇到的问题

所遇问题皆于以上提出并给出了解决方案。

特别鸣谢

[老鸽](#)



[How](#)

[士+](#)

[老e](#)

[郑老师](#)